

Parallel Computing - MPI

Message Passing Interface



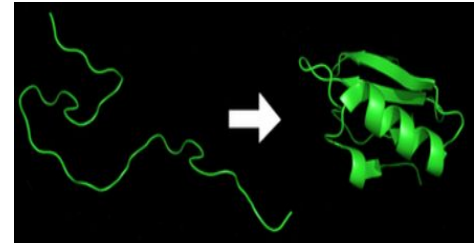
Samir Shaikh
HPC - Tech, CDAC Pune

Agenda

- **Why Parallel Computing ?**
- **Why we need ever-increasing Performance ?**
- **Parallel programming Architectures/Model ..**
- **MPI - Message Passing Interface**
 - **What is MPI ?, Need and Evolution of MPI.**
 - **MPI program - Compile and Execution**
 - **MPI Program Structure**
 - **MPI Routines**
 - **.....**

Why we need Ever-Increasing Performance ?

- Accurate medical imaging
- Fast and accurate web searches
- Realistic computer games, Entertainment
- Climate modeling
- Protein folding
- Artificial Intelligence
- Energy research
- Data analysis
-

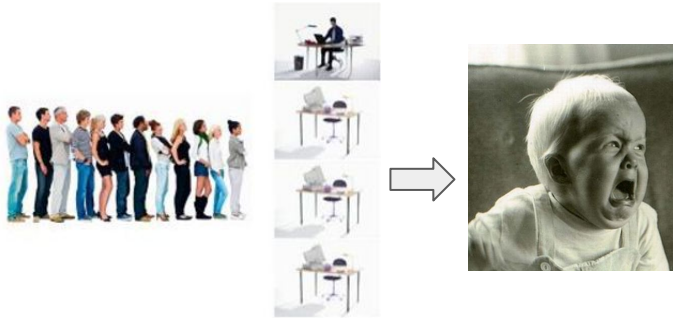


Why Parallel Computing ?

- Aren't single processor systems fast enough ?

Why Parallel Computing ?

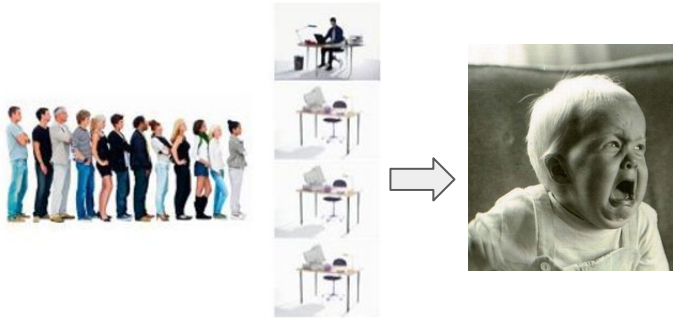
- Aren't single processor systems fast enough ?



Serial Computing

Why Parallel Computing ?

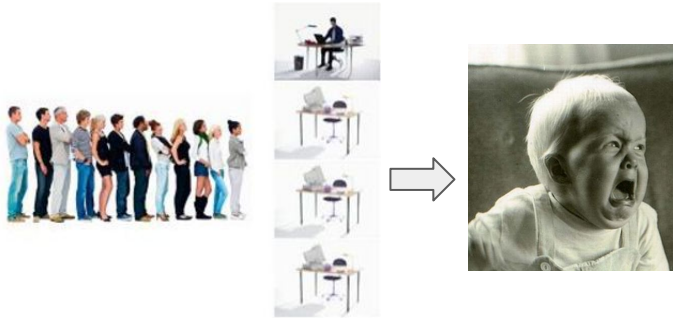
- Aren't single processor systems fast enough ?
- Why to build parallel systems ? Why build systems with multiple processors ?



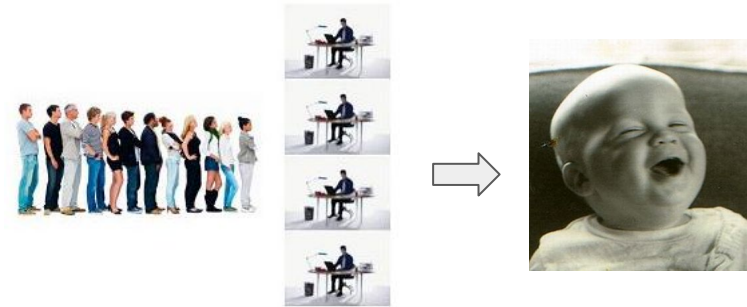
Serial Computing

Why Parallel Computing ?

- Aren't single processor systems fast enough ?
- Why to build parallel systems ? Why build systems with multiple processors ?



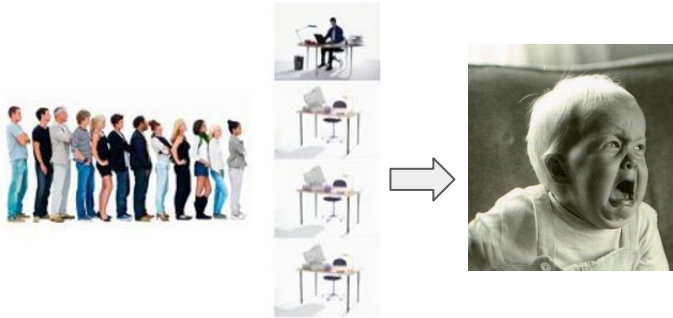
Serial Computing



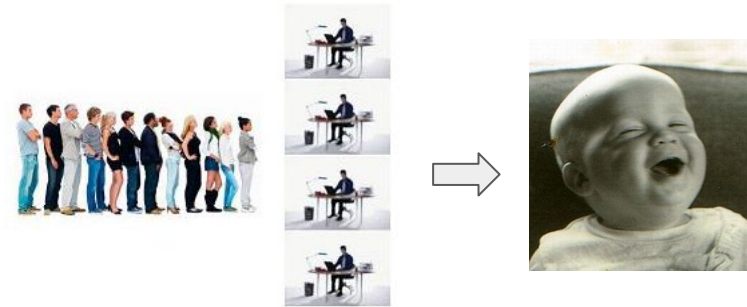
Parallel Computing

Why Parallel Computing ?

- Aren't single processor systems fast enough ?
- Why to build parallel systems ? Why build systems with multiple processors ?
- Why can't we write programs that will automatically convert serial programs to parallel programs ?



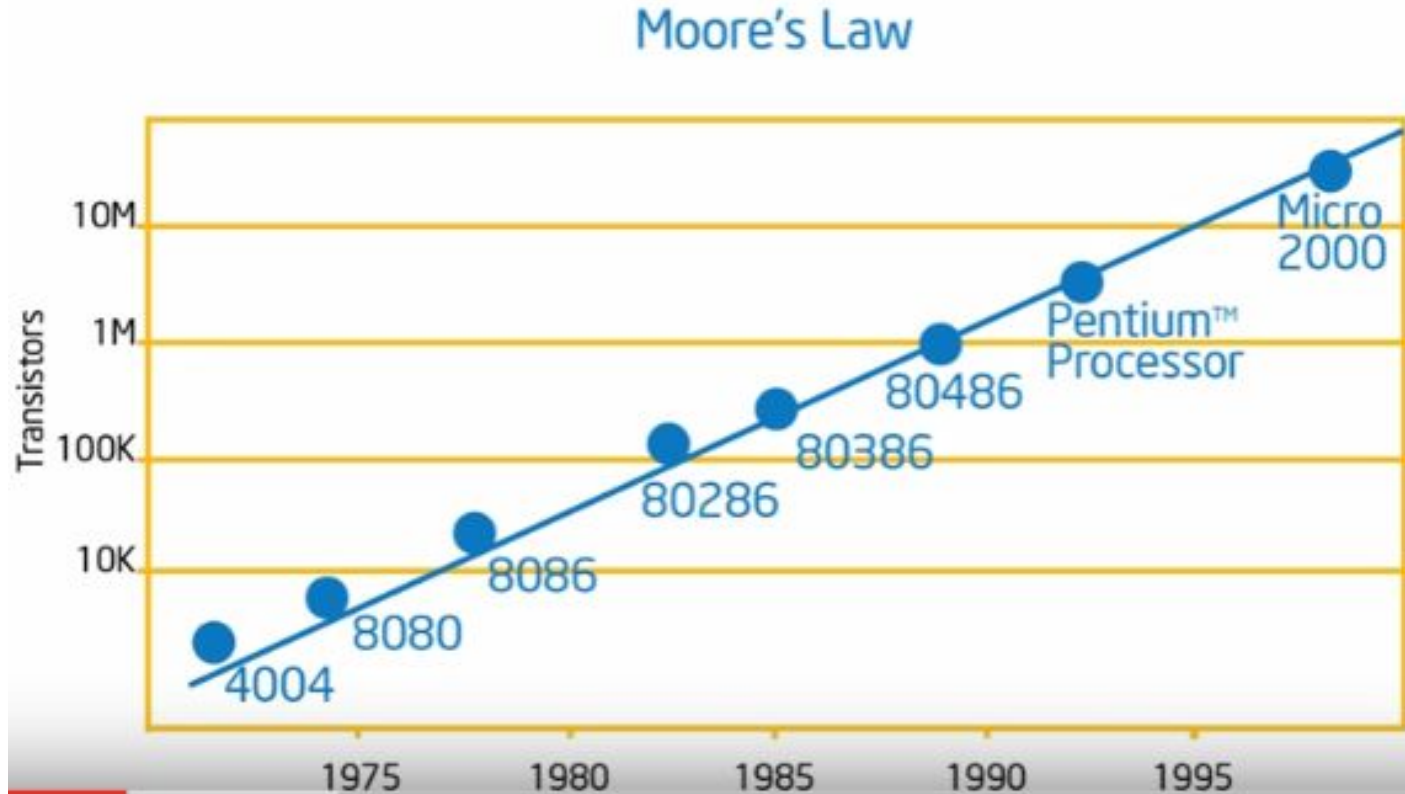
Serial Computing



Parallel Computing

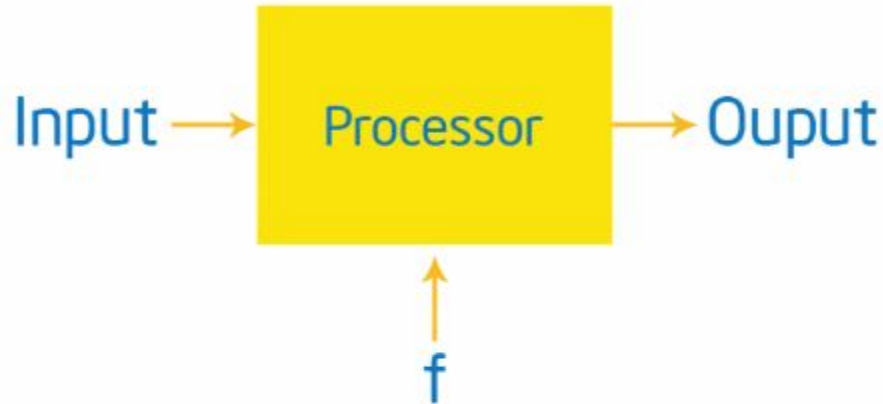
What Moore's Law tells.. ?

What Moore's Law tells.. ?



Uniprocessor ?

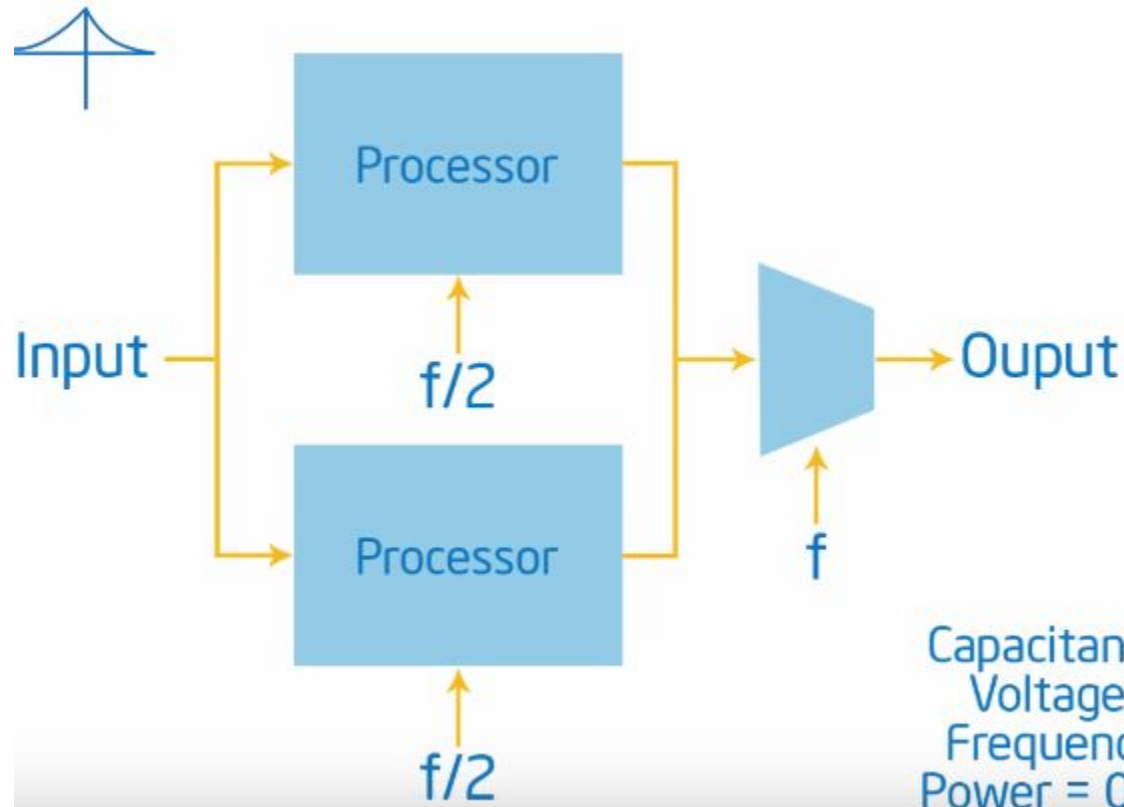
Uniprocessor ?



Capacitance = C
Voltage = V
Frequency = f
Power = CV^2f


Parallel Architecture ?

Parallel Architecture ?



Parallel Programming Models..


❑ Shared-memory Model




- **UMA - Uniform Memory Access**
- **NUMA - Non-Uniform Memory Access**

Parallel Programming Models..

❑ Shared-memory Model




- UMA - Uniform Memory Access
- NUMA - Non-Uniform Memory Access



❖ **openMP**
❖ **Pthreads ...**


Parallel Programming Models..

❑ Distributed-memory Model




Parallel Programming Models..

❑ Distributed-memory Model




❖ MPI - Message Passing Interface

MPI - Message Passing Interface




MPI - Message Passing Interface

- The Message Passing Interface Standard (MPI) is a message passing library standard based on the consensus of the MPI Forum





MPI - Message Passing Interface

- The Message Passing Interface Standard (MPI) is a message passing library standard based on the consensus of the MPI Forum
- In MPI a Message is passed from one process to another process




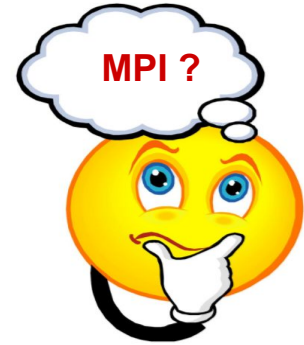
MPI - Message Passing Interface

- The Message Passing Interface Standard (MPI) is a message passing library standard based on the consensus of the MPI Forum
- In MPI a Message is passed from one process to another process





MPI - Message Passing Interface

- The Message Passing Interface Standard (MPI) is a message passing library standard based on the consensus of the MPI Forum
- In MPI a Message is passed from one process to another process
- MPI is based on Routines.



MPI - Message Passing Interface

- The Message Passing Interface Standard (MPI) is a message passing library standard based on the consensus of the MPI Forum
- In MPI a Message is passed from one process to another process
- MPI is based on Routines.
- MPI is not an IEEE or ISO standard, but has in fact, become the "industry standard" for writing message passing programs on HPC platforms.




MPI - Development

- The MPI standard has gone through a number of revisions, with the most recent version being MPI-3.x
 - MPI-3.1 - Jun 2015
 - MPI-3.0 - Sep 2012 Standard was approved
 - MPI-2.2 - Sep 2009
 - MPI-2.1 - Sep 2008
 - MPI-1.3 - May 2008
 - MPI-1.2 - July 1997
 - MPI-1.1 - June 1995
 - MPI-1.0 - May 1994

MPI - Development

- The MPI standard has gone through a number of revisions, with the most recent version being MPI-3.x
 - MPI-3.1 - Jun 2015
 - MPI-3.0 - Sep 2012 Standard was approved
 - MPI-2.2 - Sep 2009
 - MPI-2.1 - Sep 2008
 - MPI-1.3 - May 2008
 - MPI-1.2 - July 1997
 - MPI-1.1 - June 1995
 - MPI-1.0 - May 1994


Wait....
..Answer me first



- ❖ **What is Process ?**
- ❖ **Is MPI a new programming Language ..?**




The Goal ..?




How to Achieve it ..?


How to Achieve it ..?




How to Achieve it ..?




How to Achieve it ..?




How to Achieve it ..?




How to Achieve it ..?





How to Achieve it ..?




How ..?




How ..?



.....




MPI




How MPI Works ..?

How MPI Works ..?




How MPI Works ..?

- ❖ Creates Instances of same program on Every Processor involved..!




How MPI Works ..?

- ❖ Creates Instances of same program on Every Processor involved..!



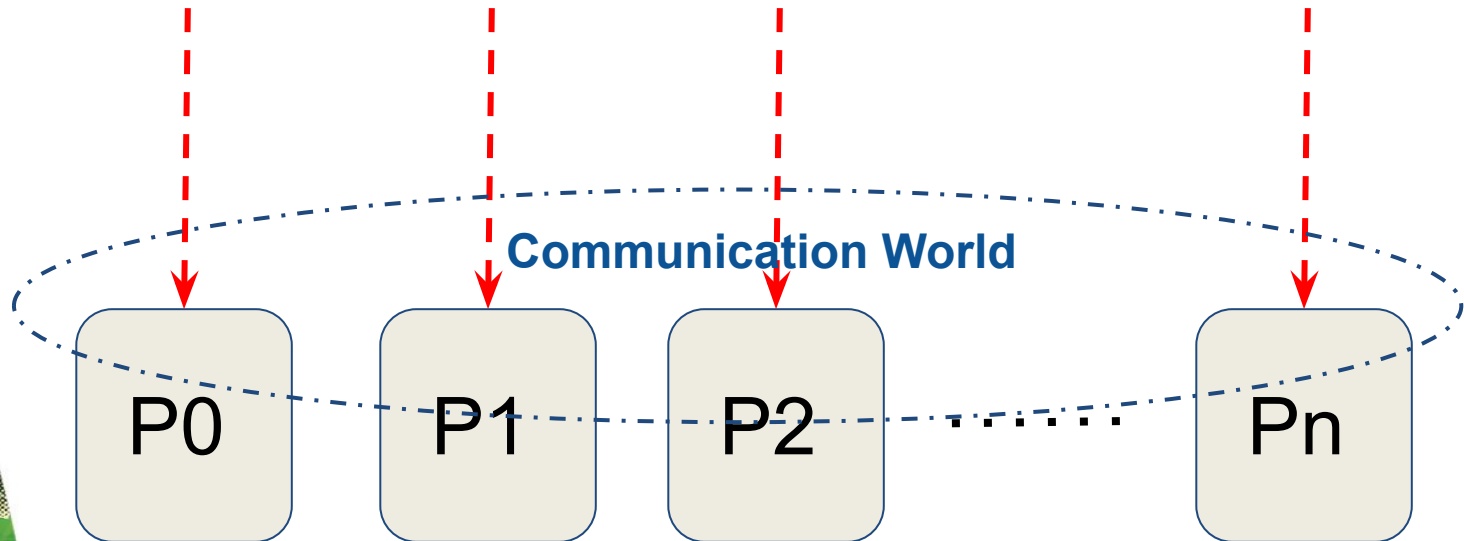
How MPI Works ..?

- ❖ Creates Instances of same program on Every Processor involved..!



How MPI Works ..?

- ❖ Creates Instances of same program on Every Processor involved..!



❖ Got it ?

❖ Got it ?



❖ Got it ?



MPI - Message Passing Interface

MPI is built on 'Routines'

The basic MPI Routines :-

- ☐ MPI_Init () ;
- ☐ MPI_Comm_rank () ;
- ☐ MPI_Comm_size () ;
- ☐ MPI_Send () ;
- ☐ MPI_Recv () ;
- ☐ MPI_Finalize () ;

- ☐ - - - - -


MPI - Communication

MPI - Communication

Point to Point Commⁿ

MPI - Communication


Point to Point Commⁿ



MPI - Communication


Point to Point Commⁿ

Collective Commⁿ




MPI - Communication

Point to Point Commⁿ




Collective Commⁿ




MPI - Communication

Point to Point Commⁿ



Collective Commⁿ



**It's Always Better to understand anything
with example.....**


It's Always Better to understand anything with example.....



...Do You Agree ?

Numerical Integration : Trapezoidal Rule

Numerical Integration : Trapezoidal Rule




Numerical Integration : Trapezoidal Rule

Trapezoid rule for integrating $\int_a^b f(x)dx$

with $h = (b - a)/n$ is

$$f(x) \approx \frac{h}{2}(f(x_0) + f(x_n)) + h \sum_{i=1}^{n-1} f(x_i)$$

where $x_i = a + ih, i = 0, 1, \dots, n$




Numerical Integration : Trapezoidal Rule

Trapezoid rule for integrating $\int_a^b f(x)dx$

with $h = (b - a)/n$ is

$$f(x) \approx \frac{h}{2}(f(x_0) + f(x_n)) + h \sum_{i=1}^{n-1} f(x_i)$$

where $x_i = a + ih, i = 0, 1, \dots, n$




Numerical Integration : Trapezoidal Rule

Trapezoid rule for integrating $\int_a^b f(x)dx$

with $h = (b - a)/n$ is

$$f(x) \approx \frac{h}{2}(f(x_0) + f(x_n)) + h \sum_{i=1}^{n-1} f(x_i)$$

where $x_i = a + ih, i = 0, 1, \dots, n$



A decorative graphic on the left side of the slide features a dark green curved shape, a green and yellow patterned area, and a stylized map of India at the bottom left.

How do you achieve it Serially .. ?

```
/* traprul_serial.c */
```

```
float f(float x) ;
```




Function which we want to integrate

```
/* trapeule_serial.c */
```

```
float f(float x) ;
```

Function which we want to integrate




```
/* traprul_serial.c */
```

```
float f(float x);
```

Function which we want to integrate

```
float Trap(float a, float b, int n, float h)
```

```
{
```

```
    float integral, x;
```

```
    int i;
```

```
    integral = (f(a) + f(b)) / 2.0;
```

```
    x = a;
```

```
    for (i=1; i<= n-1; i++)
```

```
    {
```

```
        x = x + h;
```


```
        integral = integral + f(x);
```

```
    }
```

```
    return integral*h;
```



```
}
```



```
/* traprul_serial.c */
```

```
float f(float x);
```

Function which we want to integrate

```
float Trap(float a, float b, int n, float h)
```

```
{
```

```
    float integral, x;
```

```
    int i;
```

```
    integral = (f(a) + f(b)) / 2.0;
```

```
    x = a;
```

```
    for (i=1; i<= n-1; i++)
```

```
    {
```

```
        x = x + h;
```


```
        integral = integral + f(x);
```

```
    }
```

```
    return integral*h;
```



```
}
```



```
/* traprul_serial.c */
```

```
float f(float x);
```

Function which we want to integrate

```
float Trap(float a, float b, int n, float h)
```

```
{
```

```
    float integral, x;
```

```
    int i;
```

```
    integral = (f(a) + f(b)) / 2.0;
```

```
    x = a;
```

```
    for (i=1; i<= n-1; i++)
```

```
    {
```

```
        x = x + h;
```


```
        integral = integral + f(x);
```

```
    }
```

```
    return integral*h;
```



```
}
```



```
/* traprul_serial.c */
```

```
float f(float x);
```

Function which we want to integrate

```
float Trap(float a, float b, int n, float h)
```

```
{
```

```
    float integral, x;
```

```
    int i;
```

```
    integral = (f(a) + f(b)) / 2.0;
```

```
    x = a;
```

```
    for (i=1; i<= n-1; i++)
```

```
    {
```

```
        x = x + h;
```


```
        integral = integral + f(x);
```

```
    }
```

```
    return integral*h;
```



```
}
```



```
/* traprul_serial.c */
```

```
float f(float x);
```

Function which we want to integrate

```
float Trap(float a, float b, int n, float h)
```

```
{
```

```
    float integral, x;
```

```
    int i;
```

```
    integral = (f(a) + f(b)) / 2.0;
```

```
    x = a;
```

```
    for (i=1; i<= n-1; i++)
```

```
    {
```

```
        x = x + h;
```


```
        integral = integral + f(x);
```

```
    }
```

```
    return integral*h;
```



```
}
```



```
/* traprul_serial.c */
```

```
float f(float x);
```

Function which we want to integrate

```
float Trap(float a, float b, int n, float h)
```

```
{
```

```
    float integral, x;
```

```
    int i;
```

```
    integral = (f(a) + f(b)) / 2.0;
```

```
    x = a;
```

```
    for (i=1; i<= n-1; i++)
```

```
    {
```

```
        x = x + h;
```


```
        integral = integral + f(x);
```

```
    }
```

```
    return integral*h;
```




```
}
```




A decorative graphic on the left side of the slide features a dark green curved shape, a green and yellow patterned area, and a stylized map of India at the bottom left.

How we can do it Parallely .. ?

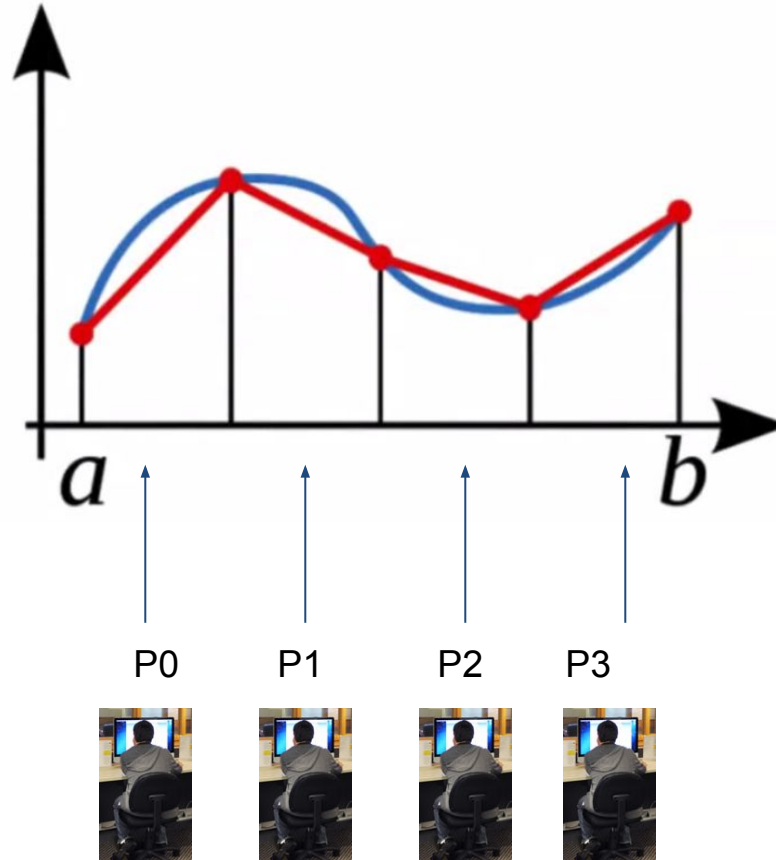
Trap Rule : Parallel Approach




Trap Rule : Parallel Approach





Trap Rule : Parallel Approach



Trap Rule : Parallel Approach



Trap Rule : Parallel Approach

 P_0 P_1 P_2 P_3 

Trap Rule : Parallel Approach

Trapezoid rule for integrating $\int_a^b f(x)dx$


with $h = (b - a)/n$ is

$$f(x) \approx \frac{h}{2}(f(x_0) + f(x_n)) + h \sum_{i=1}^{n-1} f(x_i)$$

where $x_i = a + ih, i = 0, 1, \dots, n$

Given p processes, each process can work on n/p intervals

Note: for simplicity will assume n/p is an integer



I

process	interval
0	$[a, a + \frac{n}{p}h]$
1	$[a + \frac{n}{p}h, a + 2\frac{n}{p}h]$
...	...
p-1	$[a + (p-1)\frac{n}{p}h, b]$


Trap Rule : Parallel Approach

Trapezoid rule for integrating $\int_a^b f(x)dx$

with $h = (b - a)/n$ is

$$f(x) \approx \frac{h}{2}(f(x_0) + f(x_n)) + h \sum_{i=1}^{n-1} f(x_i)$$

where $x_i = a + ih, i = 0, 1, \dots, n$



➔ Given p processes, each process can work on n/p intervals

Note: for simplicity will assume n/p is an integer

I

process	interval
0	$[a, a + \frac{n}{p}h]$
1	$[a + \frac{n}{p}h, a + 2\frac{n}{p}h]$
...	...
p-1	$[a + (p-1)\frac{n}{p}h, b]$

Trap Rule : Parallel Approach

Trapezoid rule for integrating $\int_a^b f(x)dx$


with $h = (b - a)/n$ is

$$f(x) \approx \frac{h}{2}(f(x_0) + f(x_n)) + h \sum_{i=1}^{n-1} f(x_i)$$

where $x_i = a + ih, i = 0, 1, \dots, n$

Given p processes, each process can work on n/p intervals

Note: for simplicity will assume n/p is an integer



I

process	interval
0	$[a, a + \frac{n}{p}h]$
1	$[a + \frac{n}{p}h, a + 2\frac{n}{p}h]$
...	...
p-1	$[a + (p-1)\frac{n}{p}h, b]$

Trap Rule : Parallel Approach

Trapezoid rule for integrating $\int_a^b f(x)dx$



with $h = (b - a)/n$ is

$$f(x) \approx \frac{h}{2}(f(x_0) + f(x_n)) + h \sum_{i=1}^{n-1} f(x_i)$$

where $x_i = a + ih, i = 0, 1, \dots, n$

Given p processes, each process can work on n/p intervals

Note: for simplicity will assume n/p is an integer



I

process	interval
0	$[a, a + \frac{n}{p}h]$
1	$[a + \frac{n}{p}h, a + 2\frac{n}{p}h]$
...	...
p-1	$[a + (p-1)\frac{n}{p}h, b]$

Trap Rule : Parallel Approach

Trapezoid rule for integrating $\int_a^b f(x)dx$



with $h = (b - a)/n$ is

$$f(x) \approx \frac{h}{2}(f(x_0) + f(x_n)) + h \sum_{i=1}^{n-1} f(x_i)$$

where $x_i = a + ih, i = 0, 1, \dots, n$

Given p processes, each process can work on n/p intervals

Note: for simplicity will assume n/p is an integer



I

process	interval
0	$[a, a + \frac{n}{p}h]$
1	$[a + \frac{n}{p}h, a + 2\frac{n}{p}h]$
...	...
p-1	$[a + (p-1)\frac{n}{p}h, b]$

Trap Rule : Parallel Approach

Trapezoid rule for integrating $\int_a^b f(x)dx$



with $h = (b - a)/n$ is

$$f(x) \approx \frac{h}{2}(f(x_0) + f(x_n)) + h \sum_{i=1}^{n-1} f(x_i)$$

where $x_i = a + ih, i = 0, 1, \dots, n$

Given p processes, each process can work on n/p intervals

Note: for simplicity will assume n/p is an integer



I

process	interval
0	$[a, a + \frac{n}{p}h]$
1	$[a + \frac{n}{p}h, a + 2\frac{n}{p}h]$
...	...
p-1	$[a + (p-1)\frac{n}{p}h, b]$

```
#include <stdio.h>
```

```
#include <mpi.h>
```


```
void Get_data(int p, int my_rank, double* a_p, double* b_p, int* n_p);
```

```
double Trap(double local_a, double local_b, int local_n, double h);
```

```
    /* Calculate local area */
```

```
double f(double x);
```

```
    /* function we're integrating */
```



```
#include <stdio.h>
```

```
#include <mpi.h>
```

```
void Get_data(int p, int my_rank, double* a_p, double* b_p, int* n_p);
```


```
double Trap(double local_a, double local_b, int local_n, double h);
```

```
/* Calculate local area */
```

```
double f(double x);
```

```
/* function we're integrating */
```

```
#include <stdio.h>
```



```
#include <mpi.h>
```


```
void Get_data(int p, int my_rank, double* a_p, double* b_p, int* n_p);
```

```
double Trap(double local_a, double local_b, int local_n, double h);  
/* Calculate local area */
```

```
double f(double x);  
/* function we're integrating */
```

```
#include <stdio.h>
```

```
#include <mpi.h>
```



```
void Get_data(int p, int my_rank, double* a_p, double* b_p, int* n_p);
```

```
double Trap(double local_a, double local_b, int local_n, double h);
```

```
/* Calculate local area */
```


```
double f(double x);
```

```
/* function we're integrating */
```

```
#include <stdio.h>
```

```
#include <mpi.h>
```

```
void Get_data(int p, int my_rank, double* a_p, double* b_p, int* n_p);
```



```
double Trap(double local_a, double local_b, int local_n, double h);
```

```
/* Calculate local area */
```

```
double f(double x);
```

```
/* function we're integrating */
```


```
#include <stdio.h>
```

```
#include <mpi.h>
```

```
void Get_data(int p, int my_rank, double* a_p, double* b_p, int* n_p);
```

```
double Trap(double local_a, double local_b, int local_n, double h);
```

```
/* Calculate local area */
```



```
double f(double x);
```

```
/* function we're integrating */
```



```
int main(int argc, char** argv)
{
    int      my_rank;
    int      p;
    double   a;
    double   b;
    int      n;
    double   h;
    double   local_a;
    double   local_b;
    int      local_n;
    double   my_area;
    double   total;
    int      source;
    int      dest = 0;
    int      tag = 0;
    MPI_Status status;
```


```
/* My process rank */
/* The number of processes */
/* Left endpoint */
/* Right endpoint */
/* Number of trapezoids */
/* Trapezoid base length */
/* Left endpoint my process */
/* Right endpoint my process */
/* Number of trapezoids for */
/* Integral over my interval */
/* Total area */
/* Process sending area */
/* All messages go to 0 */
```

```
MPI_Init(&argc, &argv);  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);  
MPI_Comm_size(MPI_COMM_WORLD, &p);
```

```
Get_data(p, my_rank, &a, &b, &n);
```

```
h = (b-a)/n;                                /* h is the same for all processes */  
local_n = n/p;                              /* So is the number of trapezoids */
```

```
local_a = a + my_rank*local_n*h;  
local_b = local_a + local_n*h;  
my_area = Trap(local_a, local_b, local_n, h);
```



MPI_Init(&argc, &argv);

MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

MPI_Comm_size(MPI_COMM_WORLD, &p);

Get_data(p, my_rank, &a, &b, &n);

h = (b-a)/n; **/* h is the same for all processes */**


local_n = n/p; **/* So is the number of trapezoids */**

local_a = a + my_rank*local_n*h;

local_b = local_a + local_n*h;

my_area = Trap(local_a, local_b, local_n, h);

Environment Management Routines


**`MPI_Init(&argc, &argv);`**`MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);``MPI_Comm_size(MPI_COMM_WORLD, &p);``Get_data(p, my_rank, &a, &b, &n);``h = (b-a)/n;`**`/* h is the same for all processes */`**`local_n = n/p;`**`/* So is the number of trapezoids */`**`local_a = a + my_rank*local_n*h;``local_b = local_a + local_n*h;``my_area = Trap(local_a, local_b, local_n, h);`

→ `MPI_Init(&argc, &argv);`
`MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);` 
`MPI_Comm_size(MPI_COMM_WORLD, &p);`

`Get_data(p, my_rank, &a, &b, &n);`

`h = (b-a)/n;` **`/* h is the same for all processes */`**
`local_n = n/p;` **`/* So is the number of trapezoids */`**

`local_a = a + my_rank*local_n*h;`
`local_b = local_a + local_n*h;`
`my_area = Trap(local_a, local_b, local_n, h);`



```
MPI_Init(&argc, &argv);
```

```
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
```

```
MPI_Comm_size(MPI_COMM_WORLD, &p);
```

```
Get_data(p, my_rank, &a, &b, &n);
```

```
h = (b-a)/n; /* h is the same for all processes */
```

```
local_n = n/p; /* So is the number of trapezoids */
```


```
local_a = a + my_rank*local_n*h;
```

```
local_b = local_a + local_n*h;
```

```
my_area = Trap(local_a, local_b, local_n, h);
```

```
MPI_Init(&argc, &argv);
```

```
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
```



```
MPI_Comm_size(MPI_COMM_WORLD, &p)
```



```
Get_data(p, my_rank, &a, &b, &n);
```

```
h = (b-a)/n;
```

```
/* h is the same for all processes */
```

```
local_n = n/p;
```


```
/* So is the number of trapezoids */
```

```
local_a = a + my_rank*local_n*h;
```

```
local_b = local_a + local_n*h;
```

```
my_area = Trap(local_a, local_b, local_n, h);
```

```
MPI_Init(&argc, &argv);  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);  
MPI_Comm_size(MPI_COMM_WORLD, &p);
```

 `Get_data(p, my_rank, &a, &b, &n);`


```
h = (b-a)/n;                                /* h is the same for all processes */  
local_n = n/p;                              /* So is the number of trapezoids */
```

```
local_a = a + my_rank*local_n*h;  
local_b = local_a + local_n*h;  
my_area = Trap(local_a, local_b, local_n, h);
```



```
MPI_Init(&argc, &argv);  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);  
MPI_Comm_size(MPI_COMM_WORLD, &p);
```

```
Get_data(p, my_rank, &a, &b, &n);
```




```
h = (b-a)/n;                                /* h is the same for all processes */  
local_n = n/p;                              /* So is the number of trapezoids */
```

```
local_a = a + my_rank*local_n*h;  
local_b = local_a + local_n*h;  
my_area = Trap(local_a, local_b, local_n, h);
```

```
MPI_Init(&argc, &argv);  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);  
MPI_Comm_size(MPI_COMM_WORLD, &p);
```

```
Get_data(p, my_rank, &a, &b, &n);
```




```
h = (b-a)/n;                                /* h is the same for all processes */  
local_n = n/p;                               /* So is the number of trapezoids */
```

```
local_a = a + my_rank*local_n*h;  
local_b = local_a + local_n*h;  
my_area = Trap(local_a, local_b, local_n, h);
```

```
MPI_Init(&argc, &argv);  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);  
MPI_Comm_size(MPI_COMM_WORLD, &p);
```

```
Get_data(p, my_rank, &a, &b, &n);
```

```
h = (b-a)/n;                                /* h is the same for all processes */  
local_n = n/p;                             /* So is the number of trapezoids */
```




```
local_a = a + my_rank*local_n*h;  
local_b = local_a + local_n*h;  
my_area = Trap(local_a, local_b, local_n, h);
```

```
MPI_Init(&argc, &argv);  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);  
MPI_Comm_size(MPI_COMM_WORLD, &p);
```


```
Get_data(p, my_rank, &a, &b, &n);
```

```
h = (b-a)/n;                                /* h is the same for all processes */  
local_n = n/p;                              /* So is the number of trapezoids */
```

```
local_a = a + my_rank*local_n*h;  
local_b = local_a + local_n*h;  
my_area = Trap(local_a, local_b, local_n, h);
```






```
if (my_rank == 0)
{
    total = my_area;
    for (source = 1; source < p; source++)
    {
        MPI_Recv(&my_area, 1, MPI_DOUBLE, source, tag, MPI_COMM_WORLD, &status);
        total = total + my_area;
    }
} else
{
    MPI_Send(&my_area, 1, MPI_DOUBLE, dest, tag, MPI_COMM_WORLD);
}
```



```
if (my_rank == 0)
{
    total = my_area;
    for (source = 1; source < p; source++)
    {
        MPI_Recv(&my_area, 1, MPI_DOUBLE, source, tag, MPI_COMM_WORLD, &status);
        total = total + my_area;
    }
} else
{
    MPI_Send(&my_area, 1, MPI_DOUBLE, dest, tag, MPI_COMM_WORLD);
}
```

```
if (my_rank == 0)
{
    total = my_area;
    → for (source = 1; source < p; source++)
    {
        MPI_Recv(&my_area, 1, MPI_DOUBLE, source, tag, MPI_COMM_WORLD, &status);
        total = total + my_area;
    }
} else
{
    MPI_Send(&my_area, 1, MPI_DOUBLE, dest, tag, MPI_COMM_WORLD);
}
```

```
if (my_rank == 0)
{
    total = my_area;
    for (source = 1; source < p; source++)
    {
         MPI_Recv(&my_area, 1, MPI_DOUBLE, source, tag, MPI_COMM_WORLD, &status);
        total = total + my_area;
    }
} else
{
     MPI_Send(&my_area, 1, MPI_DOUBLE, dest, tag, MPI_COMM_WORLD); 
}
```




```
if (my_rank == 0)
{
    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the area from %f to %f = %.15f\n", a, b, total);
}

MPI_Finalize();

return 0;
}

/* END of MAIN */
```




```
if (my_rank == 0)
{
    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the area from %f to %f = %.15f\n", a, b, total);
}

MPI_Finalize();

return 0;
}

/* END of MAIN */
```

```
if (my_rank == 0)
{
    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the area from %f to %f = %.15f\n", a, b, total);
}
```



MPI_Finalize();


Releases the MPI resources

```
return 0;
```

```
}
```

```
/* END of MAIN */
```


```
void Get_data(int p, int my_rank, double* a_p, double* b_p, int* n_p)
{
    int    q;
    MPI_Status status;
    if (my_rank == 0)
    {
        printf("Enter a, b, and n\n");
        scanf("%lf %lf %d", a_p, b_p, n_p);
        for (q = 1; q < p; q++) {
            MPI_Send(a_p, 1, MPI_DOUBLE, q, 0, MPI_COMM_WORLD);
            MPI_Send(b_p, 1, MPI_DOUBLE, q, 0, MPI_COMM_WORLD);
            MPI_Send(n_p, 1, MPI_INT, q, 0, MPI_COMM_WORLD);
        }
    } else
    {
        MPI_Recv(a_p, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
        MPI_Recv(b_p, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
        MPI_Recv(n_p, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
    }
}
```




```
void Get_data(int p, int my_rank, double* a_p, double* b_p, int* n_p)
{
    int    q;
    MPI_Status status;
    if (my_rank == 0)
    {
        printf("Enter a, b, and n\n");
        scanf("%lf %lf %d", a_p, b_p, n_p);
        for (q = 1; q < p; q++) {
            MPI_Send(a_p, 1, MPI_DOUBLE, q, 0, MPI_COMM_WORLD);
            MPI_Send(b_p, 1, MPI_DOUBLE, q, 0, MPI_COMM_WORLD);
            MPI_Send(n_p, 1, MPI_INT, q, 0, MPI_COMM_WORLD);
        }
    } else
    {
        MPI_Recv(a_p, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
        MPI_Recv(b_p, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
        MPI_Recv(n_p, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
    }
}
```

```
void Get_data(int p, int my_rank, double* a_p, double* b_p, int* n_p)
{
    int    q;
    MPI_Status status;
    if (my_rank == 0)
    {
        printf("Enter a, b, and n\n");
        scanf("%lf %lf %d", a_p, b_p, n_p);
        for (q = 1; q < p; q++) {
            MPI_Send(a_p, 1, MPI_DOUBLE, q, 0, MPI_COMM_WORLD);
            MPI_Send(b_p, 1, MPI_DOUBLE, q, 0, MPI_COMM_WORLD);
            MPI_Send(n_p, 1, MPI_INT, q, 0, MPI_COMM_WORLD);
        }
    } else
    {
        MPI_Recv(a_p, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
        MPI_Recv(b_p, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
        MPI_Recv(n_p, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
    }
}
```

```
void Get_data(int p, int my_rank, double* a_p, double* b_p, int* n_p)
{
    int    q;
    MPI_Status status;
    if (my_rank == 0)
    {
        printf("Enter a, b, and n\n");
        scanf("%lf %lf %d", a_p, b_p, n_p);
        for (q = 1; q < p; q++) {
            MPI_Send(a_p, 1, MPI_DOUBLE, q, 0, MPI_COMM_WORLD);
            MPI_Send(b_p, 1, MPI_DOUBLE, q, 0, MPI_COMM_WORLD);
            MPI_Send(n_p, 1, MPI_INT, q, 0, MPI_COMM_WORLD);
        }
    } else
    {
        MPI_Recv(a_p, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
        MPI_Recv(b_p, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
        MPI_Recv(n_p, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
    }
}
```



```
double Trap(double local_a, double local_b , int local_n , double h )
{
    double my_area;           /* Store my result in my_area */
    double x;
    int i;
    my_area = (f(local_a) + f(local_b))/2.0;
    x = local_a;
    for (i = 1; i <= local_n-1; i++)
    {
        x = local_a + i*h;
        my_area = my_area + f(x);
    }
    my_area = my_area*h;
    return my_area;
}
```

```
double f(double x)
{
    double return_val;

    return_val = x*x + 1.0;

    return return_val;
}

/* END of Program */
```

Recap :

Recap :

- **Serial and Parallel Approach -**

Recap :

- **Serial and Parallel Approach -**
- **MPI_Comm_rank(...)**

Recap :

- **Serial and Parallel Approach -**
- **MPI_Comm_rank(...)**
- **MPI_Comm_size(...)**

Recap :

- **Serial and Parallel Approach -**
- **MPI_Comm_rank(...)**
- **MPI_Comm_size(...)**
- **Point to point communication**

Recap :

- **Serial and Parallel Approach -**
- **MPI_Comm_rank(...)**
- **MPI_Comm_size(...)**
- **Point to point communication**
- **MPI_Send(...)**

Recap :

- **Serial and Parallel Approach -**
- **MPI_Comm_rank(...)**
- **MPI_Comm_size(...)**
- **Point to point communication**
- **MPI_Send(...)**
- **MPI_Recv(...)**

Recap :


- **Serial and Parallel Approach -**
- **MPI_Comm_rank(...)**
- **MPI_Comm_size(...)**
- **Point to point communication**
- **MPI_Send(...)**
- **MPI_Recv(...)**
- **Blocking and Non-blocking Point to Point communication - Cases !**

Recap :

- **Serial and Parallel Approach -**
- **MPI_Comm_rank(...)**
- **MPI_Comm_size(...)**
- **Point to point communication**
- **MPI_Send(...)**
- **MPI_Recv(...)**
- **Blocking and Non-blocking Point to Point communication - Cases !**
- **..... Trapezoidal Rule Example**

References :

- [1] Barker, Brandon. "Message passing interface (mpi)." *Workshop: High Performance Computing on Stampede*. Vol. 262. 2015.
- [2] Yuan, Chung-Tsz, and Shenjian Chen. "Message Passing Interface (MPI)." (1996).
- [3] <https://computing.llnl.gov/tutorials/mpi/>



MPI_Comm_rank(.....)

MPI_Comm_rank(.....)


Syntax :

➔ `MPI_Comm_rank (MPI_Comm communicator , int * rank) ;`

MPI_Comm_rank(.....)

Syntax :


➔ `MPI_Comm_rank (MPI_Comm communicator , int * rank) ;`



MPI_Comm_rank(.....)

Syntax :


➔ `MPI_Comm_rank (MPI_Comm communicator , int * rank) ;`



MPI_Comm_rank(....)

Syntax :

➔ `MPI_Comm_rank (MPI_Comm communicator , int * rank) ;`



MPI_Comm_size(.....)

MPI_Comm_size(.....)


Syntax :

➔ `MPI_Comm_size (MPI_Comm communicator , int * size) ;`

MPI_Comm_size(.....)

Syntax :


➔ `MPI_Comm_size (MPI_Comm communicator , int * size) ;`



MPI_Comm_size(.....)

Syntax :


➔ `MPI_Comm_size (MPI_Comm communicator , int * size) ;`



MPI_Comm_size(.....)

Syntax :


➔ `MPI_Comm_size (MPI_Comm communicator , int * size) ;`



MPI_Comm_size(.....)

Syntax :

➔ `MPI_Comm_size (MPI_Comm communicator , int * size) ;`



MPI_Send(.....)

MPI_Send(.....)

Syntax :

➔ **MPI_Send** (**void*** msg_buffer , **Int** msg_size, **MPI_Datatype** msg_type,
Int destination, **Int** tag , **MPI_Comm** communicator) ;

MPI_Send(.....)

Syntax :

→ MPI_Send (void* ¹msg_buffer , Int ²msg_size, MPI_Datatype ³msg_type,
Int ⁴destination, Int ⁵tag , MPI_Comm ⁶communicator) ;

MPI_Send(.....)

Syntax :

→ MPI_Send (¹void* msg_buffer , ²Int msg_size, ³MPI_Datatype msg_type,
⁴Int destination, ⁵Int tag , ⁶MPI_Comm communicator) ;

¹ Address of Message buffer

MPI_Send(.....)

Syntax :

➔ **MPI_Send** (**void*** ¹msg_buffer , **Int** ²msg_size, **MPI_Datatype** ³msg_type,
Int ⁴destination, **Int** ⁵tag , **MPI_Comm** ⁶communicator) ;

¹ Address of Message buffer

² Message size

MPI_Send(.....)

Syntax :

→ MPI_Send (¹void* msg_buffer , ²Int msg_size, ³MPI_Datatype msg_type,
⁴Int destination, ⁵Int tag , ⁶MPI_Comm communicator) ;

- ¹ Address of Message buffer
- ² Message size
- ³ Data Type

MPI_Send(.....)

Syntax :

➔ **MPI_Send** (**void*** **1** **msg_buffer** , **Int** **2** **msg_size**, **MPI_Datatype** **3** **msg_type**,
Int **4** **destination**, **Int** **5** **tag** , **MPI_Comm** **6** **communicator**) ;

1 Address of Message buffer

2 Message size

3 Data Type

4 Destination process rank

MPI_Send(.....)

Syntax :

➔ **MPI_Send** (**void*** ¹msg_buffer , **Int** ²msg_size, **MPI_Datatype** ³msg_type,
Int ⁴destination, **Int** ⁵tag , **MPI_Comm** ⁶communicator) ;

- | | | | |
|--------------|---------------------------|--------------|-----------------------------|
| ¹ | Address of Message buffer | ⁴ | Destination process rank |
| ² | Message size | ⁵ | Tag - Message Identifier/.. |
| ³ | Data Type | | |

MPI_Send(.....)

Syntax :

➔ **MPI_Send** (**void*** ¹msg_buffer , **Int** ²msg_size, **MPI_Datatype** ³msg_type,
Int ⁴destination, **Int** ⁵tag , **MPI_Comm** ⁶communicator) ;

- | | | | |
|--------------|---------------------------|--------------|-----------------------------|
| ¹ | Address of Message buffer | ⁴ | Destination process rank |
| ² | Message size | ⁵ | Tag - Message Identifier/.. |
| ³ | Data Type | ⁶ | Communicator |

MPI_Send(.....)

Syntax :

➔ MPI_Send (void* ¹msg_buffer , Int ²msg_size, MPI_Datatype ³msg_type,
Int ⁴destination, Int ⁵tag , MPI_Comm ⁶communicator) ;

MPI_Send(.....)

Syntax :

➔ **MPI_Send** (**void*** ¹ **msg_buffer** , **Int** ² **msg_size** , **MPI_Datatype** ³ **msg_type** ,
Int **destination** , **Int** ⁴ **tag** , **MPI_Comm** ⁵ **communicator**) ;
⁶

1

2


3

MPI_Send(.....)

Syntax :

→ MPI_Send (¹void* msg_buffer , ²Int msg_size, ³MPI_Datatype msg_type, ⁴Int destination, ⁵Int tag , ⁶MPI_Comm communicator) ;

Tell us information about message



MPI_Send(.....)

Syntax :

→ MPI_Send (¹void* ²msg_buffer , ²Int ²msg_size, ³MPI_Datatype ³msg_type, ⁴Int ⁴destination, ⁵Int ⁵tag , ⁶MPI_Comm ⁶communicator) ;

Tell us information about message



MPI_Send(.....)

Syntax :

→ MPI_Send (¹void* msg_buffer , ²Int msg_size, ³MPI_Datatype msg_type, ⁴Int destination, ⁵Int tag , ⁶MPI_Comm communicator) ;

Tell us information about message



¹ ² ³

Tell us, where and How to send a message



⁴ ⁵ ⁶

MPI_Recv(.....)

MPI_Recv(.....)

Syntax :

➔ **MPI_Recv** (**void*** msg_buffer , **Int** buf_size, **MPI_Datatype** buf_type,
Int source, **Int** tag , **MPI_Comm** communicator, **MPI_Status***);

MPI_Recv(.....)

Syntax :

→ MPI_Recv (void* ¹msg_buffer , Int ²buf_size, MPI_Datatype ³buf_type,
Int ⁴source, Int ⁵tag , MPI_Comm ⁶communicator, MPI_Status* ⁷);

MPI_Recv(.....)

Syntax :

→ MPI_Recv (void* ¹msg_buffer , Int ²buf_size, MPI_Datatype ³buf_type,
Int ⁴source, Int ⁵tag , MPI_Comm ⁶communicator, MPI_Status* ⁷);

- ¹ Address of Message buffer
- ² Buffer size
- ³ Data Type

MPI_Recv(.....)

Syntax :

→ MPI_Recv (**1** void* **2** msg_buffer , **3** Int buf_size, MPI_Datatype buf_type, **4** Int source, **5** Int tag , MPI_Comm **6** communicator, **7** MPI_Status*);

- | | | | |
|----------|---------------------------|----------|-----------------------------|
| 1 | Address of Message buffer | 4 | Source process rank |
| 2 | Buffer size | 5 | Tag - Message Identifier/.. |
| 3 | Data Type | 6 | Communicator |

MPI_Recv(.....)

Syntax :

→ **MPI_Recv** (**void*** **msg_buffer** , **Int** **buf_size**, **MPI_Datatype** **buf_type**,
Int **source**, **Int** **tag** , **MPI_Comm** **communicator**, **MPI_Status***);

1 2 3
4 5 6 7

- | | | | |
|---|---------------------------|---|-----------------------------|
| 1 | Address of Message buffer | 4 | Source process rank |
| 2 | Buffer size | 5 | Tag - Message Identifier/.. |
| 3 | Data Type | 6 | Communicator |
| | | 7 | Status of Received message |

Successful transmission of Message..

- ➔ **MPI_Send** (**void*** msg_buffer , **Int** msg_size, **MPI_Datatype** msg_type, **Int** destination, **Int** tag , **MPI_Comm** communicator) ;
- ➔ **MPI_Recv** (**void*** msg_buffer , **Int** buf_size, **MPI_Datatype** buf_type, **Int** source, **Int** tag , **MPI_Comm** communicator, **MPI_Status***);

Successful transmission of Message..

➔ **MPI_Send** (**void*** msg_buffer , **Int** msg_size, **MPI_Datatype** msg_type, **Int** destination, **Int** tag , **MPI_Comm** communicator) ;

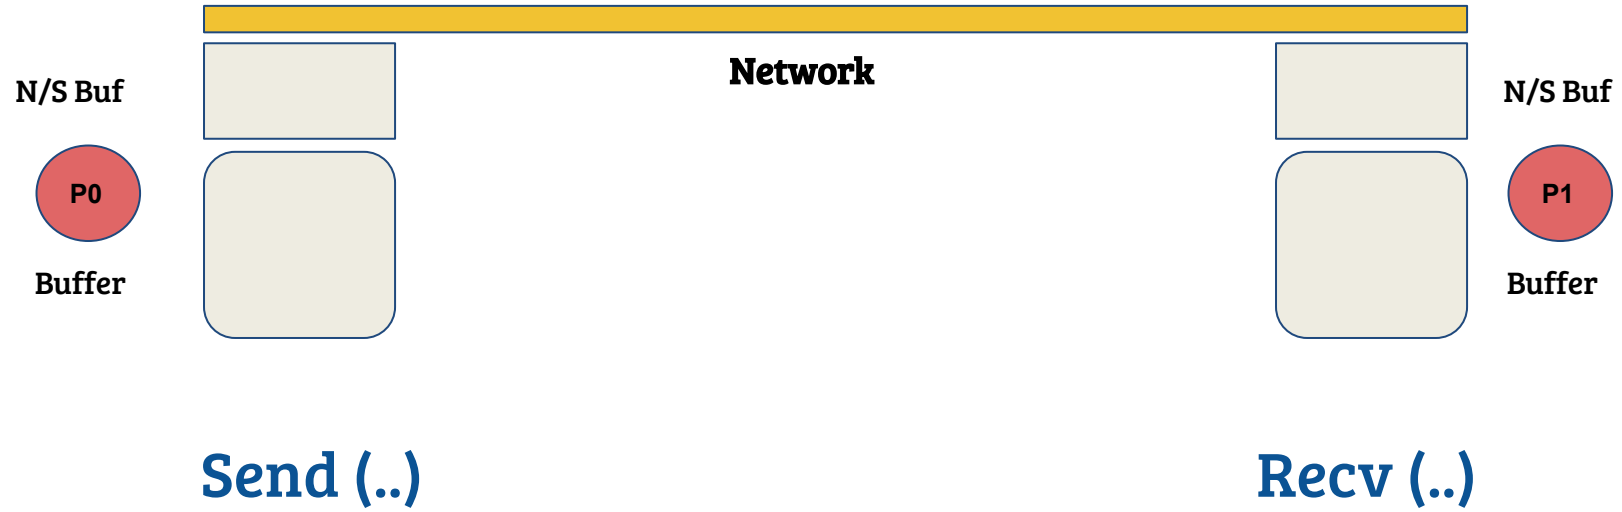
➔ **MPI_Recv** (**void*** msg_buffer , **Int** buf_size, **MPI_Datatype** buf_type, **Int** source, **Int** tag , **MPI_Comm** communicator, **MPI_Status***);

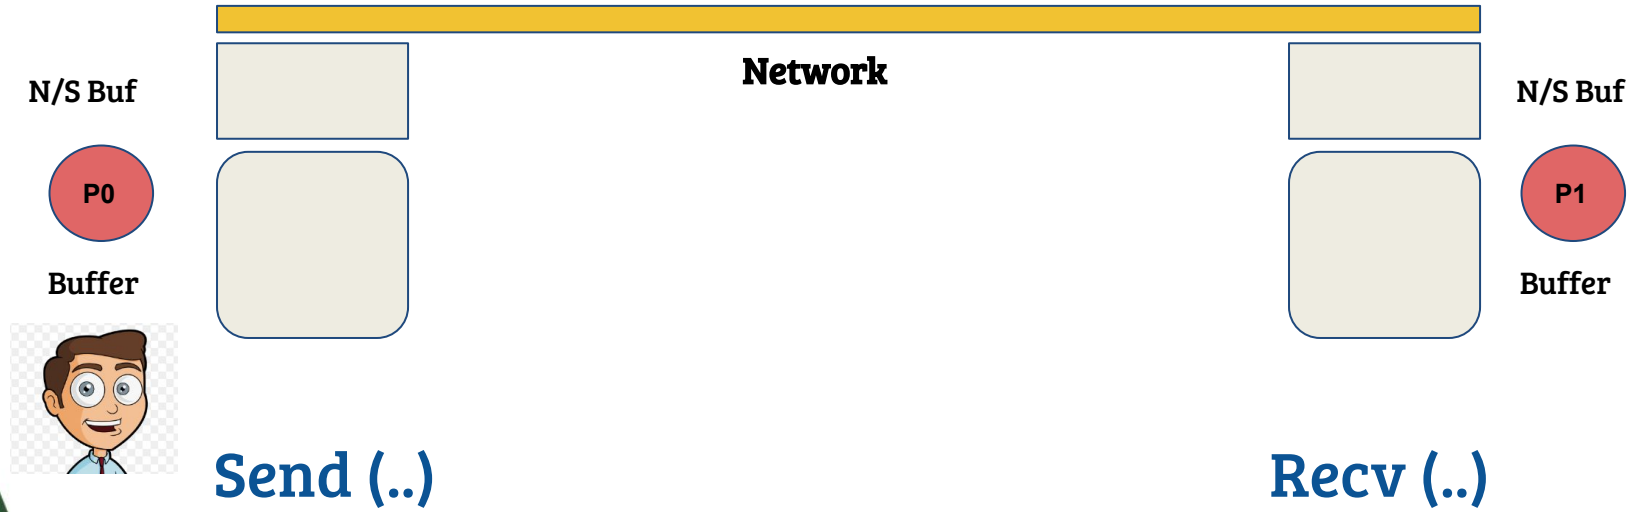
```
recv_comm = send_comm  
recv_tag  = send_tag  
dest = Destination process rank  
Src = Source process rank
```




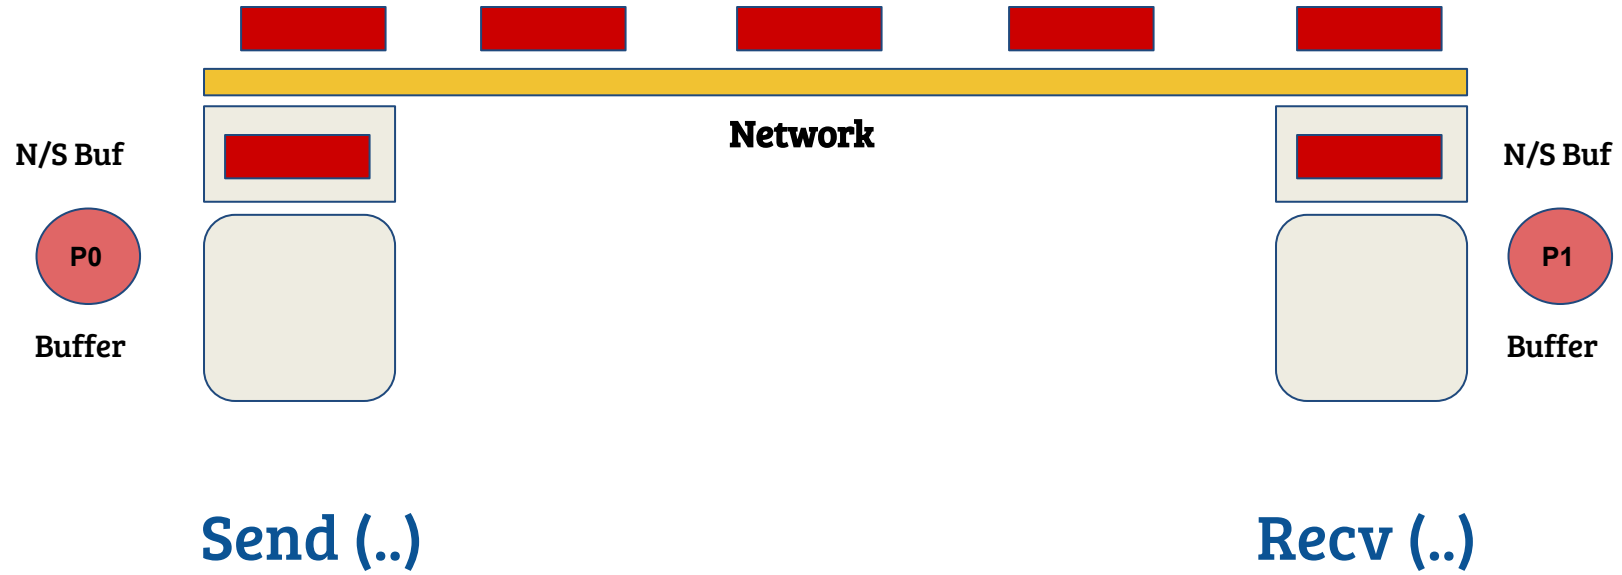
How message is transferred ... ?

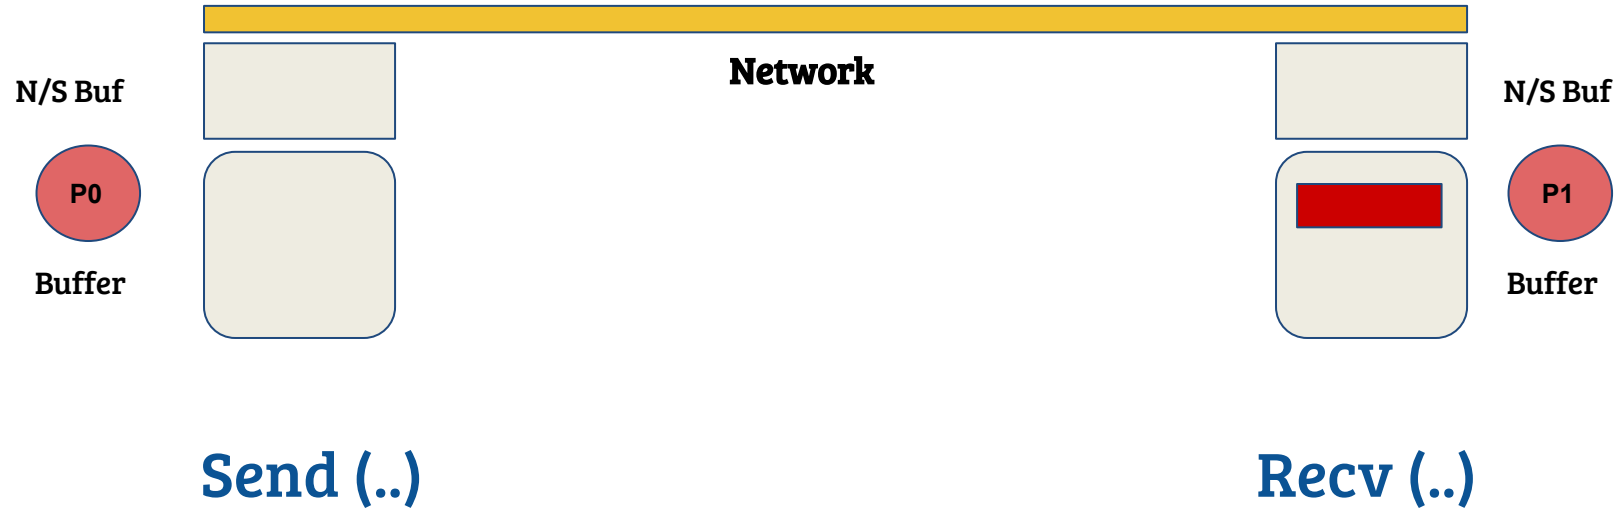
... Different cases !














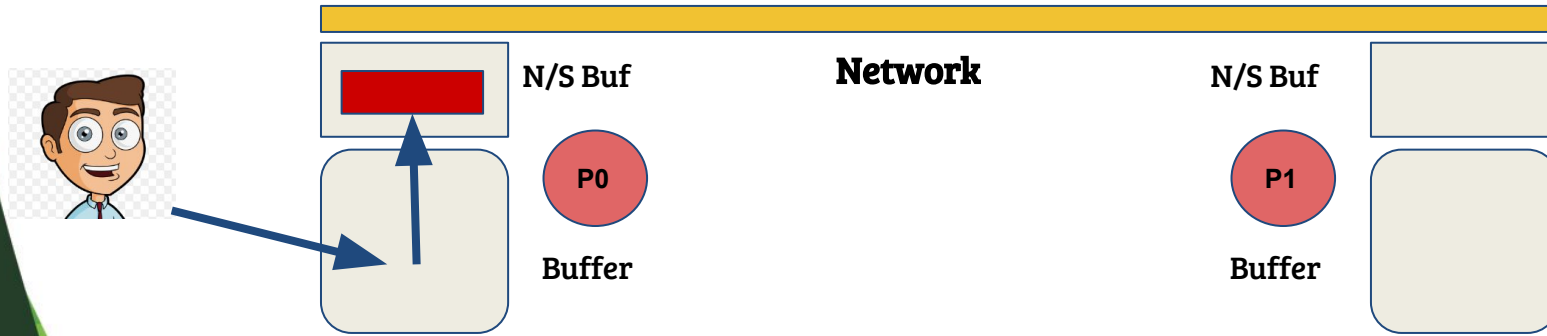




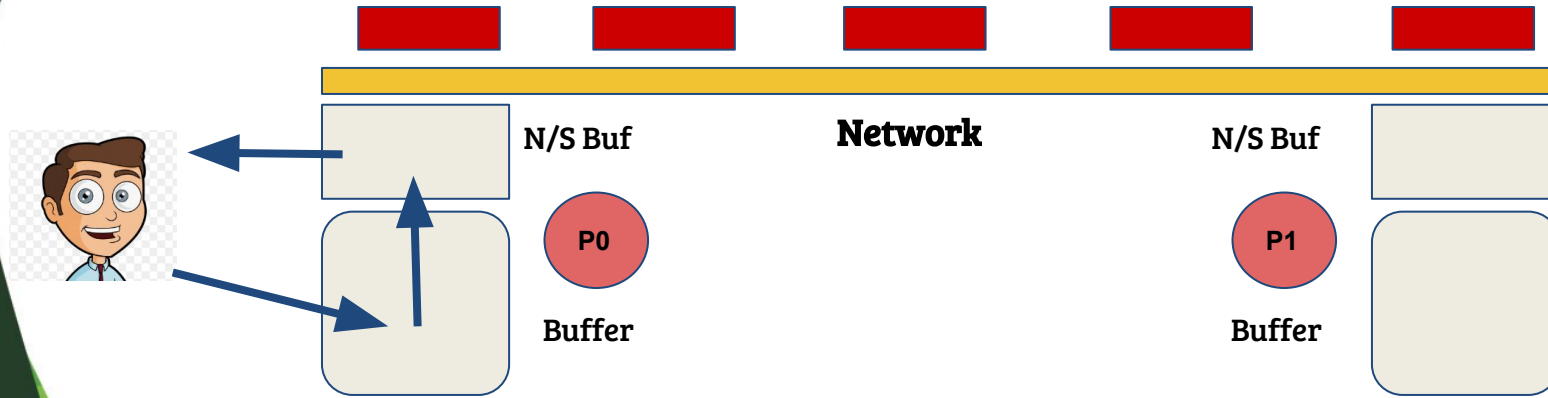
Case 1 : Blocking - Point to Point Communication



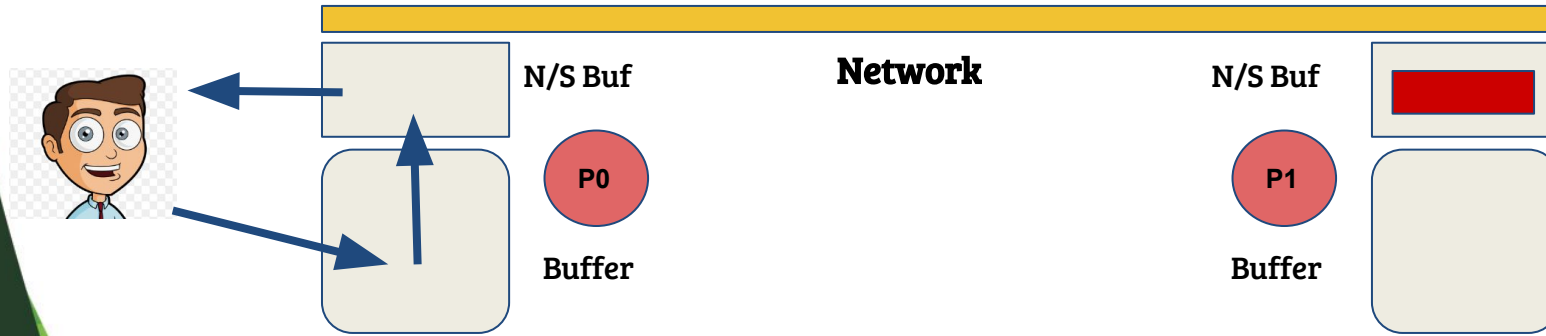
Case 1 : Blocking - Point to Point Communication



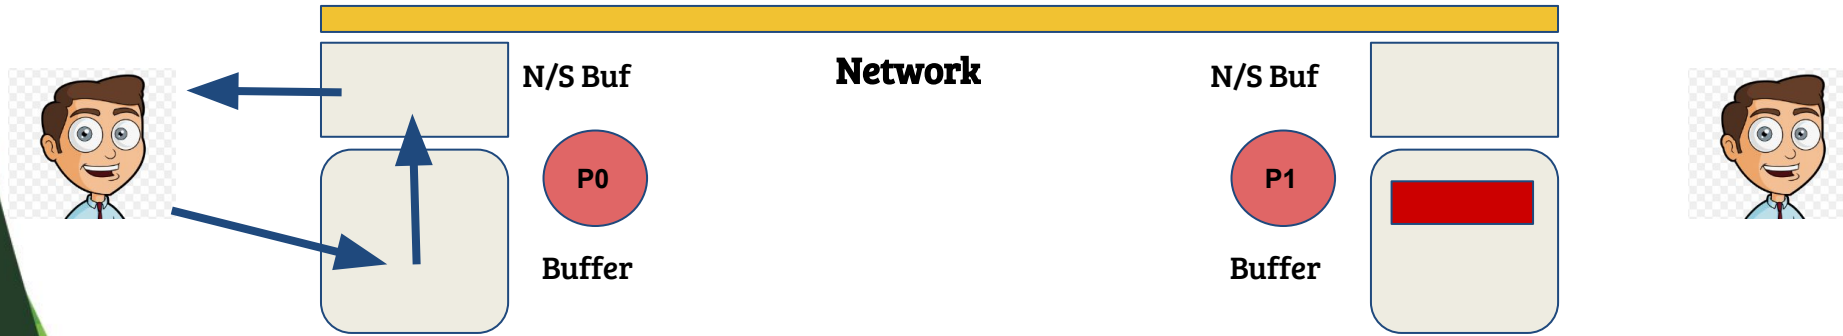
Case 1 : Blocking - Point to Point Communication




Case 1 : Blocking - Point to Point Communication



Case 1 : Blocking - Point to Point Communication




Case 1 : Blocking - Point to Point Communication




Case 1 : Blocking - Point to Point Communication


- ➡ **MPI_Send** (void* msg_buffer , Int msg_size, MPI_Datatype msg_type, Int destination, Int tag , MPI_Comm communicator) ;
- ➡ **MPI_Recv** (void* msg_buffer , Int buf_size, MPI_Datatype buf_type, Int source, Int tag , MPI_Comm communicator, MPI_Status*);



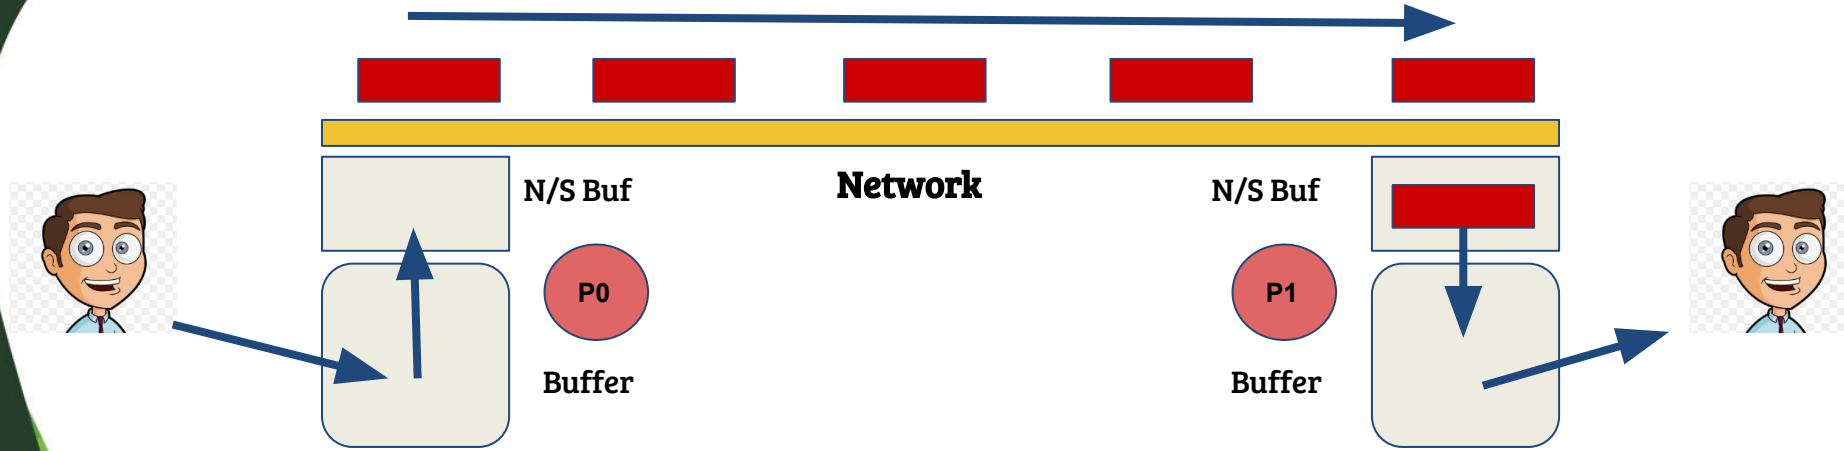
Case 2 : Synchronous Blocking - Point to Point Communication




Case 2 : Synchronous Blocking - Point to Point Communication




Case 2 : Synchronous Blocking - Point to Point Communication



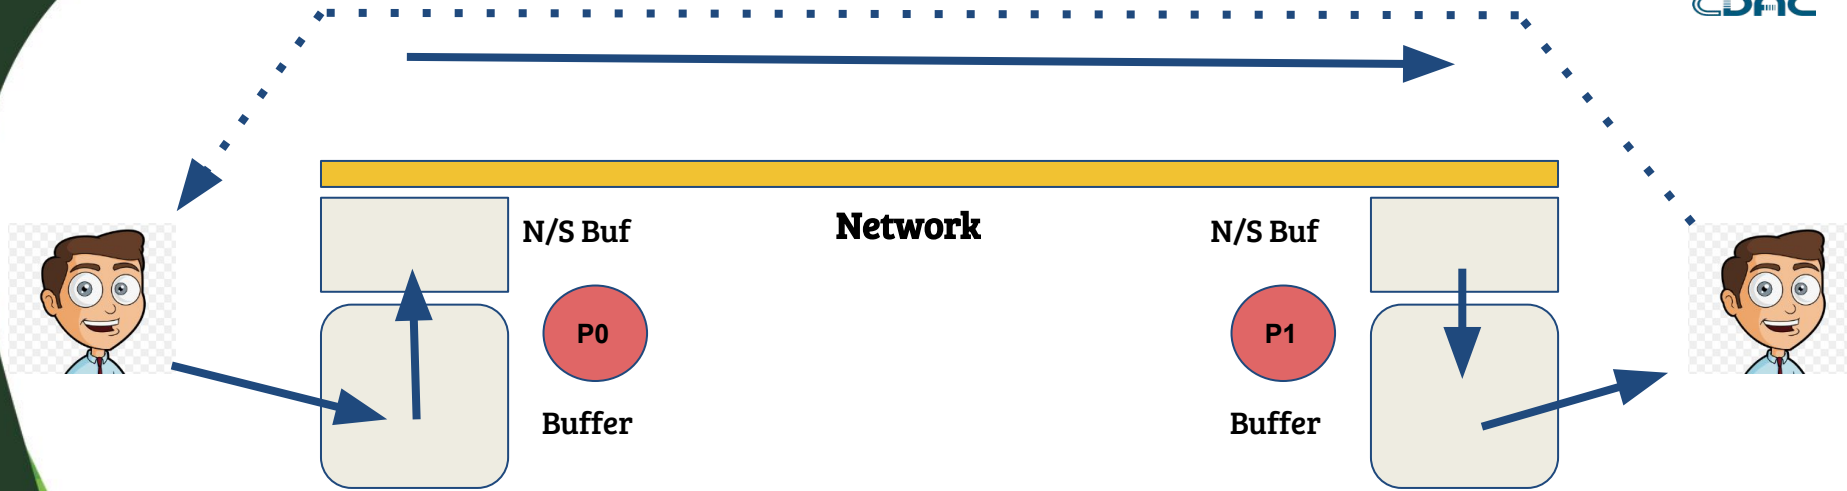
Case 2 : Synchronous Blocking - Point to Point Communication



Case 2 : Synchronous Blocking - Point to Point Communication




Case 2 : Synchronous Blocking - Point to Point Communication




Case 2 : Synchronous Blocking - Point to Point Communication


➡ **MPI_Ssend** (void* msg_buffer , Int msg_size, MPI_Datatype msg_type,
Int destination, Int tag , MPI_Comm communicator) ;



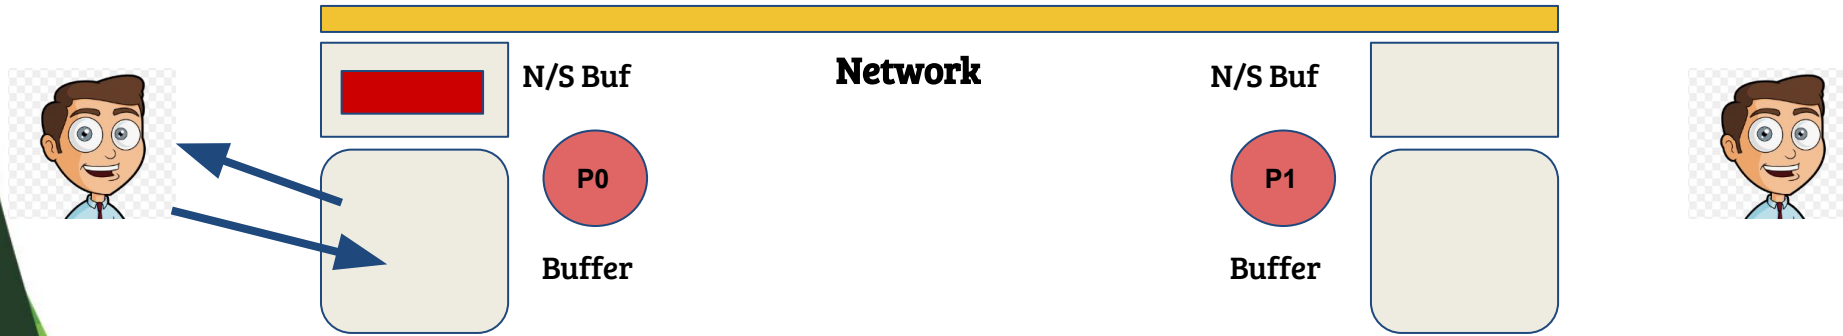
★ Case 3 : Non Blocking - Point to Point Communication




★ Case 3 : Non Blocking - Point to Point Communication



★ Case 3 : Non Blocking - Point to Point Communication




★ Case 3 : Non Blocking - Point to Point Communication




★ Case 3 : Non Blocking - Point to Point Communication

➡ `MPI_Isend (void* msg_buffer , Int msg_size, MPI_Datatype msg_type,
Int destination, Int tag , MPI_Comm communicator, req *) ;`




★ Case 3 : Non Blocking - Point to Point Communication



★ Case 3 : Non Blocking - Point to Point Communication

➡ How we will know, whether that message has been transferred or not ?




★ Case 3 : Non Blocking - Point to Point Communication

➡ How we will know, whether that message has been transferred or not ?

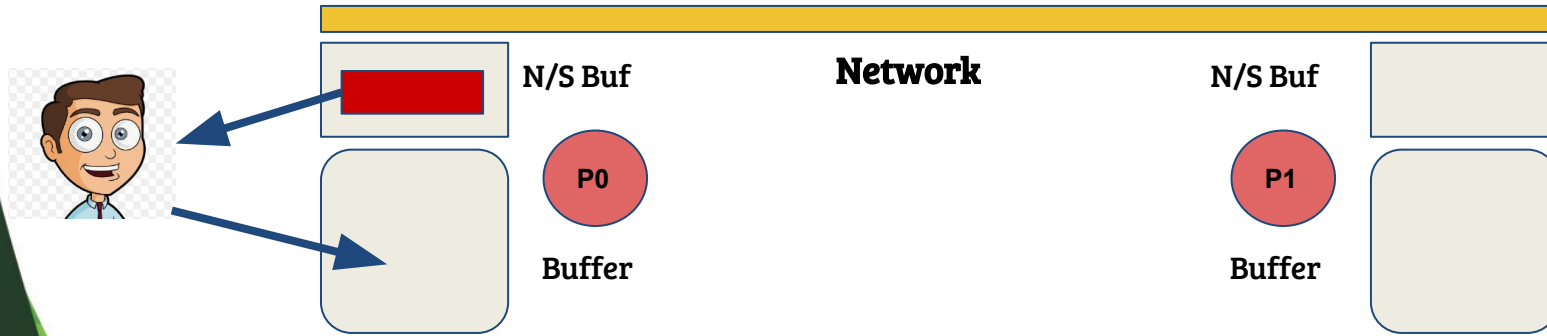


`MPI_Test(MPI_Request *req, int *flag, MPI_Status *status) ;`



★ Case 3 : Non Blocking - Point to Point Communication

➡ Data may be in buffer till next message arrive ..!




★ Case 3 : Non Blocking - Point to Point Communication

➡ Data may be in buffer till next message arrive ..!



MPI_Wait(MPI_Request *request, MPI_Status *status)



★ Case 3 : Non Blocking - Point to Point Communication




Data may be in buffer till next message arrive ..!



MPI_Wait(MPI_Request *request, MPI_Status *status)




Block and Wait till
operation get
finish..



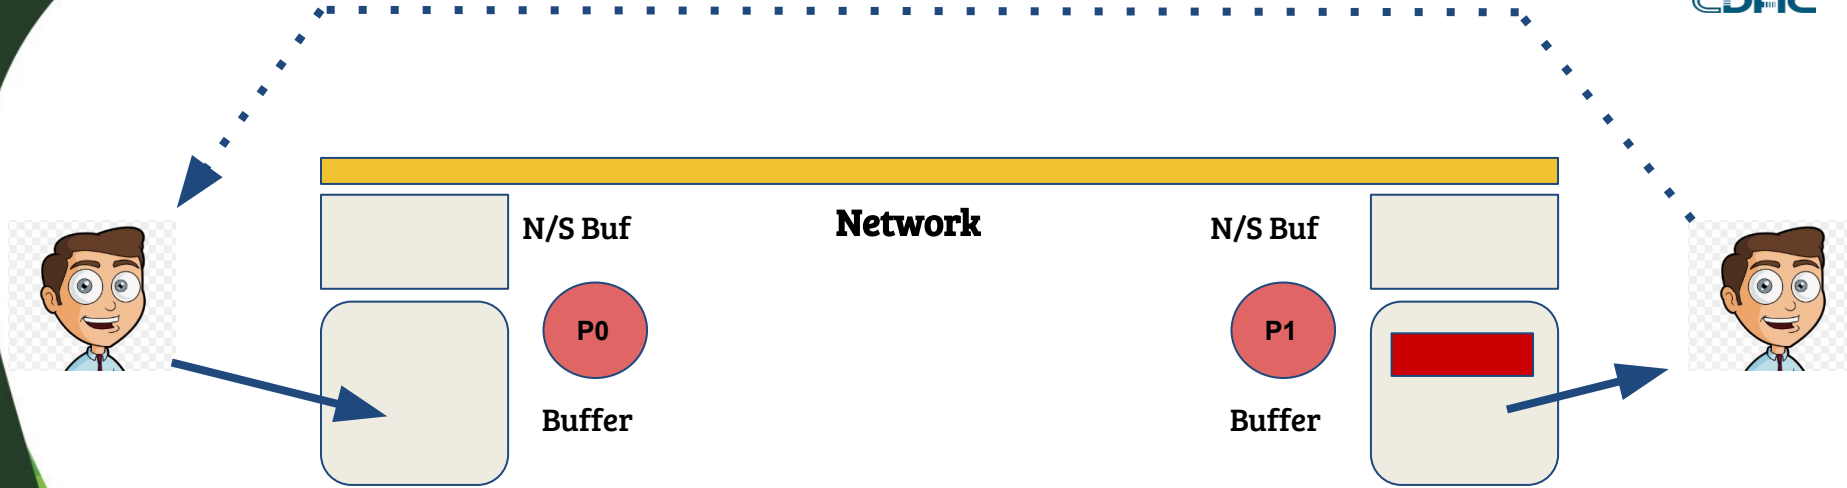
Case 4 : Synchronous Non Blocking - Point to Point Communication

➡ `MPI_issend (void* msg_buffer , Int msg_size, MPI_Datatype msg_type,
Int destination, Int tag , MPI_Comm communicator, req *) ;`



Case 4 : Synchronous Non Blocking - Point to Point Communication

➡ `MPI_issend (void* msg_buffer , Int msg_size, MPI_Datatype msg_type,
Int destination, Int tag , MPI_Comm communicator, req *) ;`



Case 4 : Synchronous Non Blocking - Point to Point Communication

➡ Similar to Non-Blocking **except** `MPI_Wait(...)` and `MPI_Test(...)`
Completes their operation when **destination process receive the message**

❖ Got it ?

❖ Got it ?



❖ Got it ?



